

Aryasomayajula Ram Bharadwaj

London, UK — [GitHub](#) — [LinkedIn](#) — [Blog](#) — ram.bharadwaj.arya@gmail.com

PROFESSIONAL SUMMARY

AI safety researcher studying evaluation awareness and AI control in frontier language models, with prior experience building production AI systems.

EXPERIENCE

AI Safety Researcher - LASR Labs @ Arcadia Impact

Jan 2026 – Present

London AI Safety Research (LASR) Labs

EvalDetectBench: A Benchmark for Measuring Evaluation Awareness in Frontier Language Models. X. Li*, K. Ochwang'i*, **Aryasomayajula R. Bharadwaj***, A. Souly, R. Kirk (*equal contribution, co-first authors). *Accepted to the NeurIPS 2026 Evaluations & Datasets Track (Poster); arXiv:2609.01611.*

- Co-designed an open, Inspect-compatible pipeline for measuring evaluation awareness in frontier LLMs.
- Identified a novel generator-identity confound in evaluation-awareness measurement: the model generating deployment transcripts systematically biases judge scores – an effect prior work did not isolate, and one large enough to reorder model rankings.
- Proposed a covariate-adjusted AUROC that corrects for the confound, restoring measurement validity with minimal error.
- Designed a framework to run red-team/blue-team control evaluations to test whether coding agents can covertly probe for environment realism, comparing various red-team strategies against blue-team monitors to see how they fare under a constrained audit budget.
- Helped secure \$125,000 in follow-on funding to extend our research on evaluation awareness.

Associate Technical Architect

Nov 2024 – Dec 2025

Quantiphi Analytics, Bengaluru

- Designed AI agent system for automated issue severity classification and escalation management.
- Architected conversational AI-agent chatbot for a major telecom company's sales team, refactoring legacy systems to LangGraph and reducing codebase size significantly.
- Revamped RAG data ingestion pipeline, improving retrieval accuracy by 20% and reducing time to first token by 3x.

Senior Developer - Innovation & Development Labs

June 2019 – Nov 2024

Mu Sigma Business Solutions, Bengaluru

- Led a team building a semi-autonomous data analysis platform using AutoGen, with automated prompt optimization and locally hosted LLMs.
- Built and maintained a high-velocity trading platform: migrated deployment to Kubernetes, rewrote trade-signal generation from R to Scala (Akka), enabling near real-time portfolio visualization.
- Developed ML model operationalization platform with automatic retraining pipelines and canary/blue-green deployment strategies.
- Received 6 Star Performer awards and 2 Impact Awards for technical leadership and delivery excellence.

FELLOWSHIPS & RESIDENCIES

AI Resident - [Lossfunk AI Residency](#)

May 2025 – June 2025

- Selected (5% acceptance, 10 researchers) for an in-person residency focused on independent technical research in interpretability.
- Built **PID-Steering**: extended activation steering by treating the steering coefficient as a controlled variable updated each token via PID feedback on an observed behavioural signal, addressing the brittleness of open-loop steering vectors.

RESEARCH & PUBLICATIONS

EvalDetectBench: A Benchmark for Measuring Evaluation Awareness in Frontier Language Models **NeurIPS 2026 ED Track (Poster)**

*X. Li**, *K. Ochwang'i**, **Aryasomayajula R. Bharadwaj***, *A. Souly*, *R. Kirk* *Equal contribution (co-first authors)

- Open, Inspect-compatible benchmark for evaluation-awareness measurement. Identifies two methodological choices in prior work – probe-question transfer across model families and transcript-generator identity – that introduce systematic bias and can reorder model rankings; proposes per-model probe calibration and a stratified generator-harmonisation correction.

Understanding Hidden Computations in Chain-of-Thought Reasoning **arXiv:2412.04537, Dec 2024**

Sole author

- Investigated whether transformer models trained with filler-token CoT (e.g., "...") perform genuinely hidden serial computation, or whether the discarded reasoning trace remains recoverable from internal activations – directly relevant to faithful-CoT and unfaithful chain-of-thought debates.
- Used layer-wise logit-lens analysis and token-rank inspection to decode the original non-filler characters from intermediate residual-stream representations without loss of downstream performance, showing the computation is encoded in the activations rather than inaccessible.

Tracing Evaluation-Awareness Emergence Through Training of OLMo-3 LessWrong, 2026 **Aryasomayajula R. Bharadwaj, R. Kirk**

- Analyzed the emergence of evaluation awareness across OLMo-3 training checkpoints, studying how awareness develops during training rather than only in fully trained models.
- Investigated measurement methodology and behavioral indicators relevant to evaluation-awareness research and frontier-model evaluations.

AWARDS & RECOGNITIONS

Honorable Mention - AI Alignment Awards

July 2023

AI Safety Research Competition

- Recognized among 118 global entries for a research proposal on “goal misgeneralization” in AI systems.

Honorable Mention - Eliciting Latent Knowledge

March 2022

Alignment Research Center

- Recognized for innovative approach to open research problem in AI safety.

Bronze Medal - Build-on-Redis Hackathon

February 2021

Redis Labs

- Developed text-to-code search tool using CodeBERT embeddings and Redis Stack for private repository indexing.

EDUCATION

Bachelor of Technology – Electronics and Communications Engineering

2015–2019

GMR Institute of Technology, Andhra Pradesh