

Aryasomayajula Ram Bharadwaj

London, UK — [GitHub](#) — [LinkedIn](#) — [Blog](#) — ram.bharadwaj.arya@gmail.com — +91-9108832338

PROFESSIONAL SUMMARY

AI safety researcher and engineer with 6+ years developing production AI systems and leading engineering teams. Experience spans interpretability and LLM evaluations. Currently researching evaluation awareness in large language models.

EXPERIENCE

AI Safety Researcher - LASR Labs @ Arcadia Impact

Jan 2026 – Present

London AI Safety Research (LASR) Labs

EvalAwareBench: A Benchmark for Measuring Evaluation Awareness in Frontier Language Models.

Aryasomayajula R. Bharadwaj, K. Ochwangi, X. Li. Under review at NeurIPS 2026.

- Co-designed the framework for measuring whether frontier LLMs distinguish evaluation from deployment contexts; built to integrate with any Inspect-compatible eval suite.
- Led the generator-identity analysis: showed that the model generating a deployment transcript systematically biases evaluation-awareness scores, and explored methods to mitigate the effect.

Associate Technical Architect

Nov 2024 – Dec 2025

Quantiphi Analytics, Bengaluru

- Designed AI agent system for automated issue severity classification and escalation management.
- Architected conversational AI-agent chatbot for a major telecom company's sales team, refactoring legacy systems to LangGraph and reducing codebase size significantly.
- Revamped RAG data ingestion pipeline, improving retrieval accuracy by 20% and reducing time to first token by 3x.

Senior Developer - Innovation & Development Labs

June 2019 – Nov 2024

Musigma Business Solutions, Bengaluru

- Led a team building a semi-autonomous data analysis platform using AutoGen, with automated prompt optimization and locally hosted LLMs.
- Built and maintained a high-velocity trading platform: migrated deployment to Kubernetes, rewrote trade-signal generation from R to Scala (Akka), enabling near real-time portfolio visualization.
- Developed ML model operationalization platform with automatic retraining pipelines and canary/blue-green deployment strategies.
- Received 6 Star Performer awards and 2 Impact Awards for technical leadership and delivery excellence.

FELLOWSHIPS & RESIDENCIES

AI Resident - [Lossfunk AI Residency](#)

May 2025 – June 2025

- Selected (5% acceptance rate) for intensive in-person AI residency program with 10 researchers, focused on independent technical research in interpretability. Built [PID-Steering](#), exploring activation steering with closed-loop control.

RESEARCH & PUBLICATIONS

EvalAwareBench: A Benchmark for Measuring Evaluation Awareness in Frontier Language Models
Under review, NeurIPS 2026

Aryasomayajula R. Bharadwaj, K. Ochwangi, X. Li

- Open benchmark measuring whether frontier LLMs distinguish evaluation from deployment contexts; integrates with any Inspect-compatible eval suite.

Understanding Hidden Computations in Transformer Language Models

August 2024

Lead Investigator

- Investigated and interpreted how filler tokens work in chain-of-thought reasoning in transformer language models.

AWARDS & RECOGNITIONS

AI Alignment Awards - Winner

July 2023

AI Safety Research Competition

- Selected among 118 global entries for winning research proposal on “goal misgeneralization” in AI systems.

Honorable Mention - Eliciting Latent Knowledge

March 2022

Alignment Research Center

- Recognized for innovative approach to open research problem in AI safety.

Bronze Medal - Build-on-Redis Hackathon

February 2021

Redis Labs

- Developed text-to-code search tool using CodeBERT embeddings and Redis Stack for private repository indexing.

EDUCATION

Bachelor of Technology – Electronics and Communications Engineering

2015–2019

GMR Institute of Technology, Andhra Pradesh