

Aryasomayajula Ram Bharadwaj

London, UK — [GitHub](#) — [LinkedIn](#) — [Blog](#) — ram.bharadwaj.arya@gmail.com — +91-9108832338

PROFESSIONAL SUMMARY

Engineer turned AI safety researcher. Built production AI systems and led engineering teams before moving into research on interpretability and LLM evaluations. Currently researching evaluation awareness in large language models.

EXPERIENCE

AI Safety Researcher - LASR Labs @ Arcadia Impact

Jan 2026 – Present

London AI Safety Research (LASR) Labs

EvalAwareBench: A Benchmark for Measuring Evaluation Awareness in Frontier Language Models.

K. Ochwangi*, X. Li, **Aryasomayajula R. Bharadwaj***. Under review at NeurIPS 2026. **Equal contribution.*

- Co-designed an open, Inspect-compatible pipeline for measuring evaluation awareness in frontier LLMs.
- Identified a novel generator-identity confound in evaluation-awareness measurement: the model generating deployment transcripts systematically biases judge scores – an effect prior work did not isolate, and one large enough to reorder model rankings.
- Proposed a covariate-adjusted AUROC that corrects for the confound, restoring measurement validity with minimal error.

Associate Technical Architect

Nov 2024 – Dec 2025

Quantiphi Analytics, Bengaluru

- Designed AI agent system for automated issue severity classification and escalation management.
- Architected conversational AI-agent chatbot for a major telecom company's sales team, refactoring legacy systems to LangGraph and reducing codebase size significantly.
- Revamped RAG data ingestion pipeline, improving retrieval accuracy by 20% and reducing time to first token by 3x.

Senior Developer - Innovation & Development Labs

June 2019 – Nov 2024

Musigma Business Solutions, Bengaluru

- Led a team building a semi-autonomous data analysis platform using AutoGen, with automated prompt optimization and locally hosted LLMs.
- Built and maintained a high-velocity trading platform: migrated deployment to Kubernetes, rewrote trade-signal generation from R to Scala (Akka), enabling near real-time portfolio visualization.
- Developed ML model operationalization platform with automatic retraining pipelines and canary/blue-green deployment strategies.
- Received 6 Star Performer awards and 2 Impact Awards for technical leadership and delivery excellence.

FELLOWSHIPS & RESIDENCIES

AI Resident - [Lossfunk AI Residency](#)

May 2025 – June 2025

- Selected (5% acceptance, 10 researchers) for an in-person residency focused on independent technical research in interpretability.
- Built [PID-Steering](#): extended activation steering by treating the steering coefficient as a controlled variable updated each token via PID feedback on an observed behavioural signal, addressing the brittleness of open-loop steering vectors.

RESEARCH & PUBLICATIONS

EvalAwareBench: A Benchmark for Measuring Evaluation Awareness in Frontier Language Models Under review, NeurIPS 2026

K. Ochwangi, X. Li, Aryasomayajula R. Bharadwaj* (*Equal contribution)*

- Open, Inspect-compatible benchmark for evaluation-awareness measurement. Identifies two methodological choices in prior work – probe-question transfer across model families and transcript-generator identity – that introduce systematic bias and can reorder model rankings; proposes per-model probe calibration and a stratified generator-harmonisation correction.

Understanding Hidden Computations in Chain-of-Thought Reasoning arXiv:2412.04537, Dec 2024

Sole author

- Investigated whether transformer models trained with filler-token CoT (e.g., “...”) perform genuinely hidden serial computation, or whether the discarded reasoning trace remains recoverable from internal activations – directly relevant to faithful-CoT and unfaithful chain-of-thought debates.
- Used layer-wise logit-lens analysis and token-rank inspection to decode the original non-filler characters from intermediate residual-stream representations without loss of downstream performance, showing the computation is encoded in the activations rather than inaccessible.

AWARDS & RECOGNITIONS

AI Alignment Awards - Winner

July 2023

AI Safety Research Competition

- Selected among 118 global entries for winning research proposal on “goal misgeneralization” in AI systems.

Honorable Mention - Eliciting Latent Knowledge

March 2022

Alignment Research Center

- Recognized for innovative approach to open research problem in AI safety.

Bronze Medal - Build-on-Redis Hackathon

February 2021

Redis Labs

- Developed text-to-code search tool using CodeBERT embeddings and Redis Stack for private repository indexing.

EDUCATION

Bachelor of Technology – Electronics and Communications Engineering

2015–2019

GMR Institute of Technology, Andhra Pradesh